

Kafka Streams

1.Terraform apply (updated vm size to 2xStandard_D4_v3):

```
MINGW64:/c/Users/yura1/Desktop/New folder/hw/test/kafka/task3/m12_kafkast... — X
Warning: Deprecated Resource

with kubernetes_namespace.confluent,
on main.tf line 114, in resource "kubernetes_namespace" "confluent":
114: resource "kubernetes_namespace" "confluent" {

Deprecated; use kubernetes_namespace_v1.

Releasing state lock. This may take a few moments...

Apply complete! Resources: 9 added, 0 changed, 0 destroyed.

Outputs:

acr_login_server = "acrdevwesteuropeopenq22.azurecr.io"
aks_api_server_url = <sensitive>
aks_kubeconfig = <sensitive>
aks_name = "aks-dev-westeuropenq22"
client_certificate = <sensitive>
resource_group_name = "rg-dev-westeuropenq22"
storage_account_name = "stdevwesteuropeopenq22"
storage_primary_access_key = <sensitive>

yura1@pc MINGW64 ~/Desktop/New folder/hw/test/kafka/task3/m12_kafkastreams_pytho
```

2.Confluent Platform pods (decreased number of replicas for zookeeper to 1, also reduced amount of allocated resources for some modules):

```
yura1@pc MINGW64 ~/Desktop/New folder/hw/test/kafka/task3/m12_kafkastreams_pytho
n_azure-master (main)
$ kubectl get pods -o wide
NAME                           NOMINATED NODE   READINESS   READY   STATUS    RESTARTS   AGE     IP
confluent-operator-8446f997d8-8n2gk   <none>        <none>      1/1    Running   0          23m    10.244.0.154
connect-0                         <none>        <none>      1/1    Running   0          13m    10.244.0.221
controlcenter-0                   <none>        <none>      1/1    Running   0          10m    10.244.1.101
elasticsearch-0                  <none>        <none>      1/1    Running   4 (12m ago) 13m    10.244.1.203
kafka-0                           <none>        <none>      1/1    Running   0          11m    10.244.1.131
kafka-1                           <none>        <none>      1/1    Running   0          11m    10.244.1.153
kafka-2                           <none>        <none>      1/1    Running   0          11m    10.244.0.148
ksqldb-0                          <none>        <none>      1/1    Running   0          10m    10.244.0.53
schema-registry-0                 <none>        <none>      1/1    Running   0          10m    10.244.1.160
zookeeper-0                       <none>        <none>      1/1    Running   0          13m    10.244.0.134
zookeeper-1                       <none>        <none>      1/1    Running   0          13m    10.244.1.115
zookeeper-2                       <none>        <none>      1/1    Running   0          13m    10.244.0.171
```

3.Control center at localhost:

The screenshot shows the Confluent Control Center Overview page. On the left, a sidebar lists navigation options: HOME, CONTROLCENTER.CLUSTER, Cluster overview, Brokers, Topics, Connect, ksqlDB, Consumers, Replicators, Cluster settings, and Health+. The main area is titled "Overview". It displays three main sections: "Brokers", "Topics", and "Connect".

- Brokers:** Shows 3 Total brokers, 39.31K Production (bytes / second), and 29.66K Consumption (bytes / second).
- Topics:** Shows 60 Total topics, 689 Partitions, 0 Under replicated partitions, and 0 Out-of-sync replicas.
- Connect:** Shows 1 Clusters, 0 Running, 0 Paused, 0 Degraded, and 0 Failed.

4.Kafka topic creation:

```
PS C:\Users\yural> Start-Process powershell -WindowStyle Hidden -ArgumentList 'kubectl port-forward connect-0 8083:8083
*> $null'
PS C:\Users\yural> kubectl exec kafka-0 -c kafka -- bash -c "/usr/bin/kafka-topics --create --topic expedia --replication-factor 3 --partitions 3 --bootstrap-server kafka:9092"
Created topic expedia.
PS C:\Users\yural>
```



Cluster overview

Brokers

Topics

Connect

ksqlDB

Consumers

Replicators

Cluster settings

Health+

expedia

Overview Messages Schema Configuration

General settings

name	expedia
partitions	3
min.insync.replicas	2
cleanup.policy	delete
retention.ms	604800000
max.message.bytes	1048588
retention.bytes	-1

Edit settings

Show full config

5. Data loaded to Azure Container:

Home > Resource Manager | Resource groups > rg-dev-westeurope-nq22 > stdevwesteuropenq22 | Containers >

data ...
Container

Search Overview Add Directory Upload Refresh Delete Copy Paste Rer

diagnose and solve problems Access Control (IAM) Settings Shared access tokens Manage ACL Access policy

Authentication method: Access key (Switch to Microsoft Entra user account)
Search blobs by prefix (case-sensitive)

Showing all 1 items

	Name	Last modified
<input type="checkbox"/>	topics	1/18/2026, 7:10:51 PM

6.azure-source-cc.json prepared:

```
terraform > {} azure-source-cc.json > {} config > abc azblob.account.name
1  {
2    "name": "expedia2",
3    "config": [
4      "connector.class": "io.confluent.connect.azure.blob.storage.AzureBlobStorageSourceConnector",
5      "tasks.max": "2",
6
7      "azblob.account.name": "stdevwesteuropenq22",
8      "azblob.account.key": "REDACTED",
9      "azblob.container.name": "data",
10
11     "format.class": "io.confluent.connect.azure.blob.storage.format.avro.AvroFormat",
12
13     "bootstrap.servers": "kafka:9071",
14     "topics": "expedia",
15     "topics.dir": "topics",
16
17     "transforms": "maskDate",
18     "transforms.maskDate.type": "org.apache.kafka.connect.transforms.MaskField$Value",
19     "transforms.maskDate.fields": "date_time",
20     "transforms.maskDate.replacement": "0000-00-00 00:00:00"
21   ]
22 }
```

7.Connector execution:

```
yural@pc MINGW64 ~/Desktop/New folder/hw/test/kafka/task3/m12_kafkastreams_python_azure-master/terraform (main)
$ curl -s -X POST -H "Content-Type:application/json" --data @azure-source-cc.json http://localhost:8083/connectors
```

8.Expedia Topic + Data masking demonstration:

Message details

Timestamp	Offset	Partition
1/18/2026, 7:15:32 PM	208909	0

Key	Value	Headers
1	{	
2	"id": 137419,	
3	"date_time": "0000-00-00 00:00:00",	
4	"site_name": 13,	
5	"posa_container": null,	
6	"user_location_country": 46,	
7	"user_location_region": 414,	
8	"user_location_city": 7317,	
9	"orig_destination_distance": 184.8278,	
10	"user_id": 458228,	
11	"is_mobile": 0,	
12	"is_package": 0,	
13	"channel": 9,	
14	"srch_ci": "2016-10-23",	
15	"srch_co": "2016-10-24",	
16	"srch_adults_cnt": 1,	
17	"srch_children_cnt": 0,	
18	"srch_rm_cnt": 1,	
19	"srch_destination_id": 8800,	
20	"srch_destination_type_id": 1,	
21	"hotel_id": 3143916660676	
22	}	

9.Output kafka topic creation:

```
PS C:\Users\yural> kubectl exec kafka-0 -- bash -c "/usr/bin/kafka-topics --create --topic expedia_ext --replication-factor 3 --partitions 3 --bootstrap-server kafka:9092"
WARNING: Due to limitations in metric names, topics with a period ('.') or underscore ('_') could collide. To avoid issues it is best to use either, but not both.
Created topic expedia_ext.
```

10.Pushed and deployed stram-app:

```

yural@pc MINGW64 ~/Desktop/New Folder/hw/test/kafka/task3/m12_kafkastreams_python_azure-master (main)
$ docker push "acrdevwesteuropenq22.azurecr.io"/stream-app:latest
The push refers to repository [acrdevwesteuropenq22.azurecr.io/stream-app]
138e27cba350: Pushed
7e95ebb70537: Pushed
89d573bf42b3: Pushed
91c91c91f1d2: Pushed
79d5bd8a8d26: Pushed
795dbdde24d: Pushed
1633a45502fe: Pushed
4f4fb700ef54: Pushed
081ccf923272: Pushed
c9723aa529b0: Pushed
26dfe2fac1c4: Pushed
260c7d06cc92: Pushed
4e39a27c4e8c: Pushed
Latest: digest: sha256:313b1674a965df7ad86b09e0d760dfc602d82ede8c2655d96240ce99445e7bb2 size: 856

yural@pc MINGW64 ~/Desktop/New Folder/hw/test/kafka/task3/m12_kafkastreams_python_azure-master (main)
$ kubectl apply -f stream-app.yaml
deployment.apps/kstream-app created

```

11.Messages available in expedia_ext:

Timestamp	Offset	Partition	Key	Value
1768760571085	13866	1	""	{"id":2506661,"date_time":"0000-00-00 00:00:00","site_name":2,"posa_container":null,"user_
1768760571080	13865	1	""	{"id":2506622,"date_time":"0000-00-00 00:00:00","site_name":2,"posa_container":null,"user_
1768760571075	13663	2	""	{"id":2506595,"date_time":"0000-00-00 00:00:00","site_name":2,"posa_container":null,"user_
1768760571068	13662	2	""	{"id":1034301,"date_time":"0000-00-00 00:00:00","site_name":2,"posa_container":null,"user_
1768760571065	13661	2	""	{"id":2506552,"date_time":"0000-00-00 00:00:00","site_name":11,"posa_container":null,"user_

12.KSQL Create stream and table:

having trouble? Type `help` (case-insensitive) for a rundown of how things work.

```
ksql> CREATE STREAM expedia_stream (
>     id BIGINT,
>     hotel_id BIGINT,
>     stay_category VARCHAR
> ) WITH (
>     KAFKA_TOPIC='expedia_ext',
>     VALUE_FORMAT='JSON'
> );
```

Message

Stream created

```
ksql> CREATE TABLE hotels_count AS
>     SELECT
>         stay_category,
>         COUNT(hotel_id) AS hotels_amount,
>         COUNT_DISTINCT(hotel_id) AS distinct_hotels
>     FROM expedia_stream
>     GROUP BY stay_category;
```

Message

Created query with ID CTAS_HOTELS_COUNT_1

```
ksql>     SELECT * FROM hotels_count EMIT CHANGES;
```

STAY_CATEGORY	HOTELS_AMOUNT	DISTINCT_HOTELS
Erroneous data	16	16
Standard extended stay	20	20
Long stay	11	11
Standard stay	479	439
Short stay	4349	2047
Short stay	5125	2170
Long stay	13	13
Standard stay	564	509
Erroneous data	18	18
Standard extended stay	26	26

Press CTRL-C to interrupt

12.KSQL FLOW:

[Cluster overview](#)[Brokers](#)[Topics](#)[Connect](#)[ksqlDB](#)[Consumers](#)[Replicators](#)[Cluster settings](#)

ksqldb

[Editor](#) [Flow](#) [Streams](#) [Tables](#) [Persistent queries](#) [Settings](#) SearchEXPEDIA_STREAM
STREAM

CREATE-TABLE

HOTELS_COUNT
TABLE

HOTELS_COUNT

**Data structure** TABLE**Total rows****Updates /sec**[Messages](#) [Schema](#)

Filter by keyword



No new messages

The message browser shows messages that have arrived since this page was opened.

ksqldb

Editor Flow Streams Tables Persistent queries Settings

```
1 | select * from HOTELS_COUNT;
```

● Add query properties

auto.offset.reset = Latest

+Add another field

Unable to run query.

! An unknown error has occurred. Check the connection settings.

X

For me, ksql is not working in the UI. I verified that the pod is running and tried forwarding an additional port for ksql, but it didn't help. Since everything works correctly from the command line, I conclude that this is a Windows port-forwarding-related issue.

```
yura1@pc MINGW64 ~/Desktop/New folder/hw/test/kafka/task3/m12_kafkastreams_python_azure-master (main)
$ kubectl get pod ksqldb-0
NAME      READY   STATUS    RESTARTS   AGE
ksqldb-0   1/1     Running   0          36m

yura1@pc MINGW64 ~/Desktop/New folder/hw/test/kafka/task3/m12_kafkastreams_python_azure-master (main)
$ kubectl exec -it ksqldb-0 -- curl -s http://localhost:8088/info
Defaulted container "ksqldb" out of: ksqldb, config-init-container (init)
{"KsqlServerInfo":{"version":"7.8.0","kafkaClusterId":"fAuGRENMQiqYFNxr8Zg6fw","ksqlServiceId":"confluent.ksqldb_","serverStatus":"RUNNING"}}

yura1@pc MINGW64 ~/Desktop/New folder/hw/test/kafka/task3/m12_kafkastreams_python_azure-master (main)
$ kubectl port-forward ksqldb-0 8088:8088
Forwarding from 127.0.0.1:8088 -> 8088
Forwarding from [::1]:8088 -> 8088
```

Connectio string:

ksqldb

[Editor](#) [Flow](#) [Streams](#) [Tables](#) [Persistent queries](#) [Settings](#)

General

Cluster name	ksqldb
Cluster ID	ksqldb
Endpoint	http://ksqldb.confluent.svc.cluster.local:8088
Version	7.8.0 View documentation